

## ▼ Install Transformers Library


```
!pip install transformers
```

```
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import transformers
from transformers import AutoModel, BertTokenizerFast

# specify GPU
device = torch.device("cuda")
```


## ▼ Load Dataset

```
df = pd.read_csv("spamdata_v2.csv")
df.head()
```




	label	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
df.shape
```



```
(5572, 2)
```

```
# check class distribution
df['label'].value_counts(normalize = True)
```



```
0    0.865937
1    0.134063
Name: label, dtype: float64
```

## ▼ Split train dataset into train, validation and test sets

```
train_text, temp_text, train_labels, temp_labels = train_test_split(df['text'], df['label'],
                                                                    random_state=2018,
                                                                    test_size=0.3,
                                                                    stratify=df['label'])

# we will use temp_text and temp_labels to create validation and test set
val_text, test_text, val_labels, test_labels = train_test_split(temp_text, temp_labels,
                                                                random_state=2018,
                                                                test_size=0.5,
                                                                stratify=temp_labels)
```

## ▼ Import BERT Model and BERTTokenizer

```
# import BERT-base pretrained model
bert = AutoModel.from_pretrained('bert-base-uncased')

# Load the BERT tokenizer
tokenizer = BertTokenizerFast.from_pretrained('bert-base-uncased')
```



Downloading: 100%

433/433 [00:00<00:00, 1.98kB/s]

Downloading: 100%

440M/440M [00:11<00:00, 37.5MB/s]

Downloading: 100%

232k/232k [00:39<00:00, 5.82kB/s]

```
# sample data
text = ["this is a bert model tutorial", "we will fine-tune a bert model"]

# encode text
sent_id = tokenizer.batch_encode_plus(text, padding=True, return_token_type_ids=False)

# output
print(sent_id)
```




```
{'input_ids': [[101, 2023, 2003, 1037, 14324, 2944, 14924, 4818, 102, 0], [101, 2057, 2
```

## ▼ Tokenization

```
# get length of all the messages in the train set
seq_len = [len(i.split()) for i in train_text]
```

```
pd.Series(seq_len).hist(bins = 30)
```

 <matplotlib.axes.\_subplots.AxesSubplot at 0x7efee3369828>



```
max_seq_len = 25
```

```
# tokenize and encode sequences in the training set
tokens_train = tokenizer.batch_encode_plus(
    train_text.tolist(),
    max_length = max_seq_len,
    pad_to_max_length=True,
    truncation=True,
    return_token_type_ids=False
)
```

```
# tokenize and encode sequences in the validation set
tokens_val = tokenizer.batch_encode_plus(
    val_text.tolist(),
    max_length = max_seq_len,
    pad_to_max_length=True,
    truncation=True,
    return_token_type_ids=False
)
```

```
# tokenize and encode sequences in the test set
tokens_test = tokenizer.batch_encode_plus(
    test_text.tolist(),
    max_length = max_seq_len,
    pad_to_max_length=True,
    truncation=True,
    return_token_type_ids=False
)
```

## ▼ Convert Integer Sequences to Tensors

```
# for train set
train_seq = torch.tensor(tokens_train['input_ids'])
train_mask = torch.tensor(tokens_train['attention_mask'])
train_y = torch.tensor(train_labels.tolist())

# for validation set
val_seq = torch.tensor(tokens_val['input_ids'])
val_mask = torch.tensor(tokens_val['attention_mask'])
val_y = torch.tensor(val_labels.tolist())

# for test set
test_seq = torch.tensor(tokens_test['input_ids'])
test_mask = torch.tensor(tokens_test['attention_mask'])
test_y = torch.tensor(test_labels.tolist())
```

## ▼ Create DataLoaders

```
from torch.utils.data import TensorDataset, DataLoader, RandomSampler, SequentialSampler

#define a batch size
batch_size = 32

# wrap tensors
train_data = TensorDataset(train_seq, train_mask, train_y)

# sampler for sampling the data during training
train_sampler = RandomSampler(train_data)

# dataloader for train set
train_dataloader = DataLoader(train_data, sampler=train_sampler, batch_size=batch_size)

# wrap tensors
val_data = TensorDataset(val_seq, val_mask, val_y)

# sampler for sampling the data during training
val_sampler = SequentialSampler(val_data)

# dataloader for validation set
val_dataloader = DataLoader(val_data, sampler = val_sampler, batch_size=batch_size)
```

## ▼ Freeze BERT Parameters

```
# freeze all the parameters
for param in bert.parameters():
    param.requires_grad = False
```

## ▼ Define Model Architecture

```
class BERT_Arch(nn.Module):

    def __init__(self, bert):

        super(BERT_Arch, self).__init__()

        self.bert = bert

        # dropout layer
        self.dropout = nn.Dropout(0.1)

        # relu activation function
        self.relu = nn.ReLU()

        # dense layer 1
        self.fc1 = nn.Linear(768,512)

        # dense layer 2 (Output layer)
        self.fc2 = nn.Linear(512,2)

        #softmax activation function
        self.softmax = nn.LogSoftmax(dim=1)

    #define the forward pass
    def forward(self, sent_id, mask):

        #pass the inputs to the model
        _, cls_hs = self.bert(sent_id, attention_mask=mask)

        x = self.fc1(cls_hs)

        x = self.relu(x)

        x = self.dropout(x)

        # output layer
        x = self.fc2(x)

        # apply softmax activation
        x = self.softmax(x)
```

```

    return x

# pass the pre-trained BERT to our define architecture
model = BERT_Arch(bert)

# push the model to GPU
model = model.to(device)

# optimizer from hugging face transformers
from transformers import AdamW

# define the optimizer
optimizer = AdamW(model.parameters(), lr = 1e-3)

```

## ▼ Find Class Weights

```

from sklearn.utils.class_weight import compute_class_weight

#compute the class weights
class_wts = compute_class_weight('balanced', np.unique(train_labels), train_labels)

print(class_wts)

```



[0.57743559 3.72848948]

```

# convert class weights to tensor
weights= torch.tensor(class_wts,dtype=torch.float)
weights = weights.to(device)

# loss function
cross_entropy = nn.NLLLoss(weight=weights)

# number of training epochs
epochs = 10

```

## ▼ Fine-Tune BERT

```

# function to train the model
def train():

    model.train()

    total_loss, total_accuracy = 0, 0

```

```
# empty list to save model predictions
total_preds=[]

# iterate over batches
for step,batch in enumerate(train_dataloader):

    # progress update after every 50 batches.
    if step % 50 == 0 and not step == 0:
        print(' Batch {:>5,} of {:>5,}'.format(step, len(train_dataloader)))

    # push the batch to gpu
    batch = [r.to(device) for r in batch]

    sent_id, mask, labels = batch

    # clear previously calculated gradients
    model.zero_grad()

    # get model predictions for the current batch
    preds = model(sent_id, mask)

    # compute the loss between actual and predicted values
    loss = cross_entropy(preds, labels)

    # add on to the total loss
    total_loss = total_loss + loss.item()

    # backward pass to calculate the gradients
    loss.backward()

    # clip the the gradients to 1.0. It helps in preventing the exploding gradient problem
    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)

    # update parameters
    optimizer.step()

    # model predictions are stored on GPU. So, push it to CPU
    preds=preds.detach().cpu().numpy()

    # append the model predictions
    total_preds.append(preds)

# compute the training loss of the epoch
avg_loss = total_loss / len(train_dataloader)

# predictions are in the form of (no. of batches, size of batch, no. of classes).
# reshape the predictions in form of (number of samples, no. of classes)
total_preds = np.concatenate(total_preds, axis=0)

#returns the loss and predictions
```

```
    return avg_loss, total_preds

# function for evaluating the model
def evaluate():

    print("\nEvaluating...")

    # deactivate dropout layers
    model.eval()

    total_loss, total_accuracy = 0, 0

    # empty list to save the model predictions
    total_preds = []

    # iterate over batches
    for step, batch in enumerate(val_dataloader):

        # Progress update every 50 batches.
        if step % 50 == 0 and not step == 0:

            # Calculate elapsed time in minutes.
            elapsed = format_time(time.time() - t0)

            # Report progress.
            print(' Batch {:>5,} of {:>5,}'.format(step, len(val_dataloader)))

        # push the batch to gpu
        batch = [t.to(device) for t in batch]

        sent_id, mask, labels = batch

        # deactivate autograd
        with torch.no_grad():

            # model predictions
            preds = model(sent_id, mask)

            # compute the validation loss between actual and predicted values
            loss = cross_entropy(preds, labels)

            total_loss = total_loss + loss.item()

            preds = preds.detach().cpu().numpy()

            total_preds.append(preds)

    # compute the validation loss of the epoch
    avg_loss = total_loss / len(val_dataloader)

    # reshape the predictions in form of (number of samples, no. of classes)
```



```
total_preds = np.concatenate(total_preds, axis=0)
```

```
return avg_loss, total_preds
```

## ▼ Start Model Training

```
# set initial loss to infinite
```

```
best_valid_loss = float('inf')
```

```
# empty lists to store training and validation loss of each epoch
```

```
train_losses=[]
```

```
valid_losses=[]
```

```
#for each epoch
```

```
for epoch in range(epochs):
```

```
    print('\n Epoch {:} / {:}'.format(epoch + 1, epochs))
```

```
    #train model
```

```
    train_loss, _ = train()
```

```
    #evaluate model
```

```
    valid_loss, _ = evaluate()
```

```
    #save the best model
```

```
    if valid_loss < best_valid_loss:
```

```
        best_valid_loss = valid_loss
```

```
        torch.save(model.state_dict(), 'saved_weights.pt')
```

```
    # append training and validation loss
```

```
    train_losses.append(train_loss)
```

```
    valid_losses.append(valid_loss)
```

```
    print(f'\nTraining Loss: {train_loss:.3f}')
```

```
    print(f'Validation Loss: {valid_loss:.3f}')
```



```
Epoch 1 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.526  
Validation Loss: 0.656

```
Epoch 2 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.345  
Validation Loss: 0.231

```
Epoch 3 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.344  
Validation Loss: 0.194

```
Epoch 4 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.223  
Validation Loss: 0.171

```
Epoch 5 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.219  
Validation Loss: 0.178

```
Epoch 6 / 10
  Batch    50 of   122.
  Batch   100 of   122.
```

Evaluating...

Training Loss: 0.215  
Validation Loss: 0.180

```
Epoch 7 / 10
  Batch    50 of   122.
```

Batch 100 of 122.

Evaluating...

Training Loss: 0.247

Validation Loss: 0.262

Epoch 8 / 10

Batch 50 of 122.

Batch 100 of 122.

Evaluating...

Training Loss: 0.224

Validation Loss: 0.217

Epoch 9 / 10

Batch 50 of 122.

Batch 100 of 122.

## ▼ Load Saved Model

validation loss: 0.140

#load weights of best model

path = 'saved\_weights.pt'

model.load\_state\_dict(torch.load(path))



<All keys matched successfully>

Training Loss: 0.231

## ▼ Get Predictions for Test Data

# get predictions for test data

with torch.no\_grad():

preds = model(test\_seq.to(device), test\_mask.to(device))

preds = preds.detach().cpu().numpy()

# model's performance


preds = np.argmax(preds, axis = 1)

print(classification\_report(test\_y, preds))



	precision	recall	f1-score	support
0	0.99	0.98	0.98	724
1	0.88	0.92	0.90	112
accuracy			0.97	836
macro avg	0.93	0.95	0.94	836
weighted avg	0.97	0.97	0.97	836

```
# confusion matrix
pd.crosstab(test_y, preds)
```



col_0	0	1
row_0		
0	710	14
1	9	103